

# Econometrics in R's tidyverse

by Tyler Ransom, University of Oklahoma  
[@tyleransom](#)

## Basics of R

R can be thought of as a really fancy calculator

### Packages:

- R comes with a lot of functionality out-of-the-box
- Other functionality requires the user to load packages
- One-time installation: `install.packages("tidyverse")`
- Each time you open R: `library(tidyverse)`

### Commenting:

- Use `#` to make a comment
- This tells R to ignore that code  
`# My name is Tyler`

### Assignment operator:

- Use `<-` to store a calculation, e.g. `x <- 3` (" $x = 3$ ")

### Pipe operator:

- Use `%>%` to "pipe" objects  
`y <- mean(log(x))` becomes `y <- x %>% log %>% mean`
- `%<>%` pipes forward, then backwards  
`x <- mean(log(x))` is same as `x %<>% log %>% mean`

## Working with Data

R's fundamental data object is a **data frame**

Like spreadsheets, stores data in columns and rows

**tidyverse** uses **tibbles** (enhanced data frames)

```
df <- as_tibble(mtcars)
```

### Reading in data

- Many functions for reading in different types of data  
`df <- read_csv("myfile.csv")` (comma separated)  
`df <- read_fwf("myfile.dat")` (fixed-width)
- More details: see [Data Importing Cheat Sheet](#)
- **haven** package: import foreign files (e.g. SAS, Stata, ...)

### Accessing columns of data

- To reference a column in a tibble, use `$`  
`df$mpg`  
`mean(df$mpg)` will return sample avg of `mpg` variable

### Ignore missing values

- Missing values are indicated by `NA`
- Some commands won't automatically ignore `NA` values

- For these cases, use `na.rm` option  
`mean(df$mpg, na.rm=TRUE)`  
`df$mpg %>% mean(na.rm=TRUE)` (equivalent)
- Otherwise, R would say the mean is `NA`

### Removing columns and rows from a tibble

- To keep columns in a tibble, use `select()`  
`df1 <- df %>% select(mpg, disp, hp, gear, carb)`
- To keep rows in a tibble, use `filter()`  
`df1 %<>% filter(mpg>=10)`
- To remove columns, put a minus in front  
`df1 <- df %>% select(-mpg, -disp)`

### Remove missing values from a tibble

- To remove **all** rows with *any* `NA` values, use `drop_na()`  
`df1 <- df %>% drop_na()`
- Can also drop `NA`'s from particular columns:  
`df1 <- df %>% drop_na(gear, carb)`

### Creating new columns in a tibble

- To create a new column in a tibble, use `mutate()`  
`df1 %<>% mutate(mpg_squared = mpg^2)`

### Manipulating values of a variable

- To replace (i.e. recode) values of a variable:  
`df %<>% mutate(gear = replace(gear, gear==4, 99))`  
Changes all 4's in `gear` to be 99's  
`gear==4` can be any other logical condition
- To specify a series of conditions, use `%in%`  
`df %<>% mutate(hp = replace(hp, hp %in% c(110, 120), 99))`  
Changes all 110's or 120's in `hp` to be 99's

### Working with discrete variables

- Discrete variables often require special treatment
- In R, declare discrete variables as **factors**  
`df %<>% mutate(gear = as.factor(gear))`

### Other data manipulations

- See [Data Wrangling Cheat Sheet](#)

## Getting to know your data

It's important to know what's in your data by

1. Looking at summary statistics
2. Performing cross-tabulations
3. Visualizing certain variables

### Summary statistics (skimr package)

- Report quartiles, min/max, mean, sd, and `#NA`'s:  
`skim(df)`  
or  
`df %>% skim`

### Cross-tabulations

- Report frequencies of a discrete variable:  
`table(df$gear)`
- Average  $y$  by categories of a discrete  $x$  variable:  
`df %>% group_by(gear) %>% summarize(m.mpg = mean(mpg))`

### Visualization

- Often helpful to look at a histogram or line graph
- Histogram (continuous  $x$ ):  
`ggplot(df, aes(mpg)) + geom_histogram()`
- Histogram (factor  $x$ ):  
`ggplot(df, aes(x=gear)) + geom_bar()`
- Kernel density plot:  
`ggplot(df, aes(mpg)) + geom_density()`
- Simple scatter plot with linear fit:  
`ggplot(df, aes(disp, mpg)) + geom_point() + geom_smooth(method="lm")`
- More details: see [ggplot2 Cheat Sheet](#)

## Regression modeling

### Basic OLS regression

- Regression:  

```
est <- lm(mpg ~ gear + hp, data=df)
```
- Examine regression output:  

```
summary(est)  
tidy(est)  
stargazer(est,type="text")
```
- Other functional forms:  

```
est <- lm(mpg ~ gear + I(gear^2), data=df)  
est <- lm(log(mpg) ~ gear + I(gear^2), data=df)
```
- Factor variables automatically get separate intercepts

### *t*-statics and *F*-statics

- *t*-stats, *p*-values reported in regression output
- *F*-test:  

```
linearHypothesis(est,c("gear","hp"))  
tests  $H_0 : \beta_{gear} = 0, \beta_{hp} = 0$   


```
linearHypothesis(est,c("gear=5","hp=-1"))  
tests  $H_0 : \beta_{gear} = 5, \beta_{hp} = -1$ 
```


```
- Robust *F*-test (see next section):  

```
linearHypothesis(est.rob,c("gear","hp"))
```

### Robust standard errors (estimatr package)

- Correct for heteroskedasticity:  

```
est.rob <- lm_robust(mpg ~ gear + hp, data=df)  
or  
stargazer(est,se=starpred(est.rob),type="text")
```
- Correct for serial correlation:  

```
fixed.est <- est %>% coefTest(vcov=NeweyWest)  
stargazer(est,se=list(fixed.est[,2]),type="text")
```
- Correct for clustering:  

```
est.clust <- lm_robust(mpg ~ gear + hp, data=df,  
clusters=df$carb)  
or  
stargazer(est,se=starpred(est.clust),type="text")
```

## Instrumental Variables

- Let *drat* be the endogenous covariate
- Let *wt* be the instrument
- Let *qsec* and *gear* be exogenous covariates  

```
est.iv <- ivreg(mpg ~ drat + qsec + gear |  
wt + qsec + gear, data=df)
```
- Instruments come after the | symbol
- Endogenous covariates come before the | symbol
- Exogenous covariates appear on both sides of the |
- First-stage regression:  

```
est.1 <- lm(drat ~ wt + qsec + gear, data=df)  
df %>% mutate(drat.hat = est.1$fitted.values)
```
- Second-stage regression:  

```
est.2 <- lm(mpg ~ drat.hat + qsec + gear, data=df)
```
- Can also use *estimatr* for robust SEs:  

```
est.ivr <- iv_robust(mpg ~ drat + qsec + gear |  
wt + qsec + gear, data=df)
```

## Working with time series data

- Declare a time series data frame  

```
df.ts <- zoo(df, order.by=df$year)
```
- Time series line plot:  

```
ggplot(df.ts, aes(year, inf)) + geom_line()
```
- Simple AR(1) model:  

```
est <- dynlm(inf ~ L(inf,1), data=df.ts)
```
- First-differences model:  

```
est.diff <- dynlm(d(inf) ~ unem, data = df.ts)
```
- ADF test for unit root:  

```
adf.test(df1.ts$inf, k=1)
```
- ARIMA model:  

```
est.arima <- auto.arima(df.ts$inf)
```
- Plot *h*-period-ahead forecast intervals  

```
autoplot(forecast(est.arima, h=2))
```
- Extended date and time functions available in  
*lubridate* package

## Working with panel data

- Report number of units and time periods  

```
pdim(df)
```
- Pooled OLS model  

```
est.pols <- plm(lwage ~ exper + I(exper^2) +  
year, data = df, index = c("id","year"),  
model = "pooling")
```
- Random effects model  

```
est.re <- plm(lwage ~ exper + I(exper^2) +  
year, data = df, index = c("id","year"),  
model = "random")
```
- Fixed effects model  

```
est.fe <- plm(lwage ~ exper + I(exper^2) +  
year, data = df, index = c("id","year"),  
model = "within")
```
- First differences model  

```
est.fd <- plm(lwage ~ exper + I(exper^2) +  
year, data = df, index = c("id","year"),  
model = "fd")
```

## Limited dependent variable models

### Linear probability model (LPM):

- If *y* is a factor, format it as a numeric  

```
est.lpm <- lm(as.numeric(y) ~ x1 + x2, data=df)
```

### Logit and Probit:

In this case, *y* should be formatted as a factor

- Logit:  

```
est.logit <- glm(y ~ x1 + x2,  
family=binomial(link="logit"),data=df)
```
- Probit:  

```
est.probit <- glm(y ~ x1 + x2,  
family=binomial(link="probit"),data=df)
```

## List of packages

The document requires the following packages:

tidyverse	car	zoo	forecast
magrittr	estimatr	dynlm	plm
stargazer	lmtest	AER	
broom	clubSandwich	tseries	
skimr	sandwich	lubridate	

Layout by Winston Chang, <http://wch.github.io/latexsheet/>